

JATIN PAL

+91 97298 24212 | jatinpal300003@gmail.com | linkedin.com/in/jatin-pal-ba41a828a | github.com/gilgamish-hub
Portfolio and live demos: gilgamish-hub.github.io

SUMMARY

Final-year B.Tech Computer Science student specialising in Artificial Intelligence. I build machine learning and LLM applications in Python: transformer fine-tuning, TF-IDF recommendation, and LLM features on Gemini and local Llama models. Looking for an AI/ML or Generative AI engineering role; graduating in 2027.

EDUCATION

Bennett University, Greater Noida

2023 – 2027

B.Tech in Computer Science and Engineering (Specialisation in Artificial Intelligence) | CGPA: 7.99 / 10

Coursework: Data Structures, Algorithms, Operating Systems, DBMS, Machine Learning, Generative AI

TECHNICAL SKILLS

Languages: Python, C++, Java, SQL

Machine Learning: PyTorch, TensorFlow, Keras, Scikit-learn, CNNs, GANs, VAEs, Reinforcement Learning (DQN)

Generative AI and NLP: Hugging Face Transformers, BERT, LegalBERT, DistilBERT, LangChain, LangGraph, RAG, ChromaDB, Gemini API, Ollama

Computer Vision: OpenCV, YOLOv8, U-Net, image segmentation

Backend and Data: FastAPI, REST APIs, MySQL, MongoDB, Pandas, NumPy, Matplotlib, Tableau

Tools: Git, GitHub, Docker, Docker Compose, Streamlit, Pytest, Jupyter, CUDA

PROJECTS

Legal Contract Auditor

github.com/gilgamish-hub/legal-auditor-ai

Python, PyTorch, Hugging Face Transformers, Scikit-learn, Streamlit

- Built a dataset of 9,149 clauses in 19 clause types plus "Other" from the lawyer annotations of the CUAD corpus (510 contracts), split into train, validation and test sets by contract.
- Fine-tuned LegalBERT to 88.4% accuracy and 0.864 macro F1 on 76 held-out contracts, against 81.8% and 0.811 for a TF-IDF and logistic regression baseline.
- Built a Streamlit app that reads a contract PDF, splits it into clauses, classifies each one, assigns a risk level, highlights the words that most affect the prediction, and exports a PDF audit report.

InternAI: Internship Recommendation System (Smart India Hackathon)

github.com/gilgamish-hub/innovace

Python, Scikit-learn, Gemini API, Ollama, Streamlit

- Led the team that built a content-based recommender over 6,500+ internship listings using TF-IDF and cosine similarity, with filters for location, stipend and duration.
- Added LLM features (skill-gap analysis, feedback on an uploaded resume PDF, learning roadmap) using Gemini, with automatic fallback to a local Llama 3 model through Ollama.

IT Support Agent (rule-based policy engine)

github.com/gilgamish-hub/Veridian-Corp-

Python, Pydantic, Pytest, Streamlit

- Designed a 5-stage pipeline for IT support requests (intent and entity parsing, knowledge-base and past-ticket retrieval, policy decision, tool execution, response generation) with Pydantic state models.
- Implemented 10 tool functions (account unlock, VPN renewal, laptop dispatch, security alerts, audit-log export) and a 17-test Pytest suite covering 15 request scenarios.

Genetic Algorithm Feature Selection

github.com/gilgamish-hub/genetic-feature-selection

Python, Scikit-learn, joblib, Streamlit

- Implemented a two-step feature selector: a mutual-information filter, then a genetic algorithm (tournament selection, crossover, adaptive mutation) whose fitness trades 5-fold cross-validated accuracy against feature count.
- Evaluated population fitness in parallel with joblib and compared logistic regression, random forest and SVM on the selected features in a Streamlit app.

EXPERIENCE

Data Science Intern, Codec Technologies (Remote)

Aug 2025 – Sep 2025

- Built and evaluated machine learning models in Python using Scikit-learn and Pandas.

CERTIFICATIONS

Deep Learning Specialization (DeepLearning.AI) | Getting Started with Accelerated Computing in CUDA C++ (NVIDIA)

Algorithmic Toolbox (UC San Diego) | Object-Oriented Data Structures in C++ (University of Illinois)